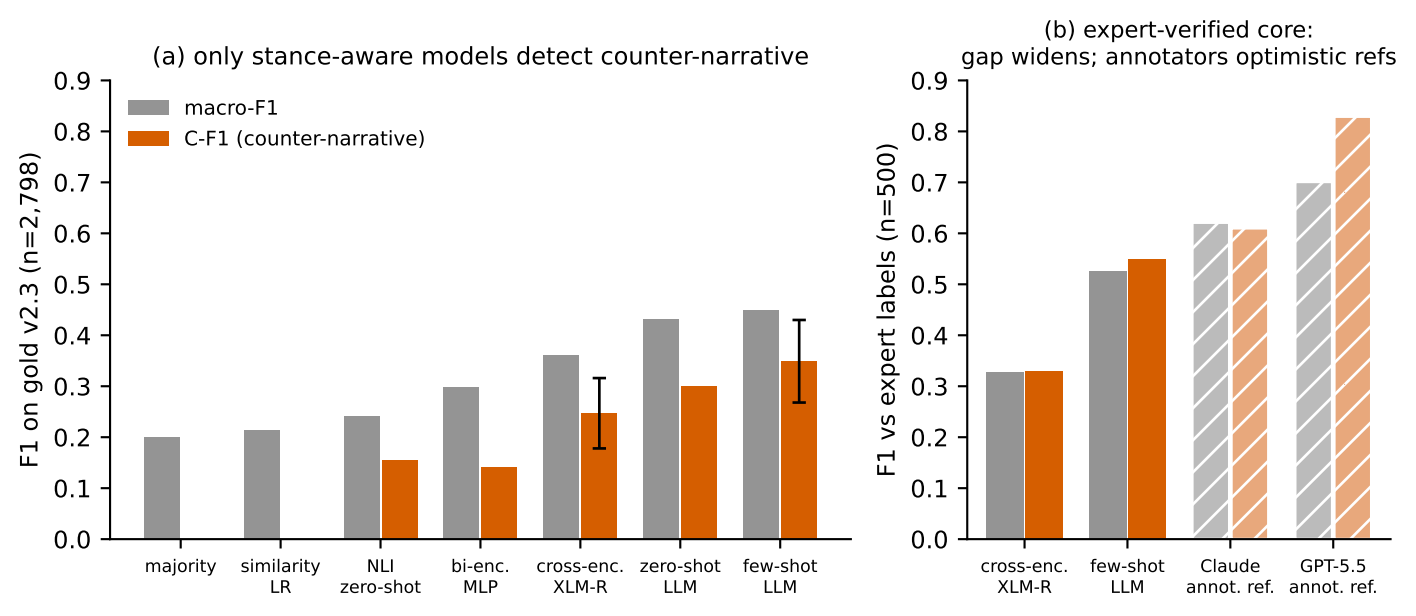
(a) only stance-aware models detect counter-narrative
macro-F1
C-F1 (counter-narrative)
F1 on gold v2.3 (n=2,798)
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
0.8
0.9
majority
similarity LR
NLI zero-shot
bi-enc. MLP
cross-enc. XLM-R
zero-shot LLM
few-shot LLM
(b) expert-verified core:
gap widens; annotators optimistic refs
F1 vs expert labels (n=500)
cross-enc. XLM-R
few-shot LLM
Claude annot. ref.
GPT-5.5 annot. ref.